

# PGA-UNet: A Lightweight Prompt-Guided Attention U-Net for Interactive Bone X-Ray Lesion Segmentation

THONG LUC<sup>1,2</sup>, HUU-BINH NGUYEN<sup>1,2</sup>, LY QUOC NGOC<sup>1,2</sup>

<sup>1</sup>B.Sc. Program in Computer Science, Faculty of Information Technology, University of Science, VNU-HCM, Vietnam

<sup>2</sup>Viet Nam National University, Ho Chi Minh City, Vietnam

Corresponding authors: Ly Quoc Ngoc (e-mail: lqngoc@fit.hcmus.edu.vn) and Thong Luc (e-mail: 22120196@hcmus.edu.vn).

**ABSTRACT** Bone X-ray lesion segmentation is difficult because lesions are often small, low-contrast, and easily confused with the surrounding anatomy. In practice, the decisive step is usually localizing the suspicious region correctly, not refining the lesion's boundary. Building on this observation, we narrow the search space before segmentation through an interactive mechanism: a clinician draws a coarse bounding box around the suspicious region, and the model uses this box as a spatial guidance signal.

This article proposes PGA-UNet, a lightweight prompt-guided CNN. The bounding box is converted into a Gaussian-smoothed plateau heatmap, then used to modulate encoder features through a Prompt Spatial Gate (PSG) and to sustain its influence in the decoder through a Conditional Attention Decoder (CAD). This design does not only use the prompt to indicate where segmentation should occur, but also keeps small-lesion-relevant features active throughout the network.

We use Attention U-Net without a prompt as the automatic reference model, and SAM-Med2D as the prompt-matched reference. Experiments on the BTXRD and FracAtlas datasets, at three resolutions (128, 256, and 512), under two box conditions, covering and controlled off-center, show that prompt-free segmentation is a very hard problem: Attention U-Net reaches only 0.52 and 0.34 Dice on the two datasets. Given the same box, PGA-UNet consistently outperforms two conventional prompt-matched baselines and fine-tuned SAM-Med2D. The margin is widest on FracAtlas, the harder dataset, and on the smallest-lesion group: fine-tuned SAM-Med2D drops to 0.26 to 0.37 Dice, while PGA-UNet still holds 0.65 to 0.72. The results are stable across repeated random data splits, and PGA-UNet is about 92 times lighter than SAM-Med2D in parameter count. In addition, a training-free self-assessment score, based on mask decisiveness and stability under prompt jitter, ranks prediction reliability without ground truth, reaching a rank correlation of about 0.5 to 0.7 with the true Dice.

It should be noted that every box in the experiments is simulated around an already-annotated lesion, so the results above only reflect segmentation ability once a good prompt is already available. This article does not answer questions about: fully automatic lesion detection, partial-coverage or off-target boxes, calibrated confidence, real clinical usability, patient-level generalization, or transfer to another dataset.

**INDEX TERMS** bone X-ray, interactive segmentation, medical image segmentation, prompt-guided learning, Attention U-Net

## I. INTRODUCTION

**T**HANKS to its low cost, speed, and accessibility, bone X-ray imaging is widely used in screening, diagnostic support, and treatment monitoring. However, segmenting lesions in this type of image is not simple when the lesion is very small, low-contrast, or partially obscured by the surrounding bone structure. In such cases, the main difficulty lies not in drawing an accurate boundary, but in localization: the model

must find the suspicious region in the whole image before it can refine any boundary.

This is why small-lesion segmentation is rarely just a boundary-drawing problem. When a lesion occupies only a tiny fraction of the image, searching the whole image is both costly and error-prone. The natural approach is to narrow the search space first, then segment within the narrowed region. Fully automatic segmentation methods, typified by U-Net and

Attention U-Net [3], [4], try to learn both localization and segmentation implicitly within the same network, with no external help at all.

Another approach is to draw on external guidance about the lesion's location. In clinical practice, a clinician can easily and quickly draw a coarse bounding box around a suspicious region, something an automatic model struggles to do well for very small lesions. Seen this way, the bounding box is not merely an interactive convenience, but a practical way to inject spatial knowledge into the model before detailed segmentation begins. Many prompt-based interactive segmentation methods have grown out of this idea, from simple operations such as clicks and scribbles [5], [6] to recent foundation models such as SAM and SAM-Med2D [1], [2].

Even so, one question remains unanswered: once the prompt has identified the region of interest, should it only act as a positional constraint for generating the mask, or should it continue to shape feature extraction, so that an already faint lesion signal is not lost through the downsampling layers? This question matters most for very small lesions, where even a small spatial deviation can substantially affect the final outcome.

This article answers that question with PGA-UNet, a lightweight prompt-guided CNN in which the bounding box is treated as a structured spatial prior, not merely an auxiliary input channel. The box is encoded into a Gaussian-smoothed plateau heatmap, used to modulate encoder features through a Prompt Spatial Gate (PSG), and then reused in the decoder through a Conditional Attention Decoder (CAD). As a result, prompt-guided lesion features stay active throughout the network, rather than being confined to limiting segmentation within the given box.

This perspective also determines how we design and read the experiments. Attention U-Net receives no spatial guidance from a clinician, so its role is only to show how hard the problem becomes when localization must be fully automatic. SAM-Med2D, by contrast, is chosen as the direct comparison because it receives the same bounding-box input. The central question of this article is therefore not whether a prompt is useful in general, but whether a lightweight architecture can exploit a clinician's coarse prompt more effectively, by preserving weak lesion signals instead of treating the prompt merely as a positional constraint.

The main contributions of this article are as follows.

- We propose PGA-UNet: a lightweight, prompt-guided architecture for interactive bone X-ray lesion segmentation.
- We build a prompt-conditioning mechanism with three parts: a Gaussian-plateau prompt, an encoder-side Prompt Spatial Gate (PSG), and a decoder-side Conditional Attention Decoder (CAD).
- We build an offline training and online interactive inference framework, together with explicit mathematical definitions for prompt encoding, encoder-side gating, and decoder-side conditioning.

- We evaluate the method on the BTXRD and FracAtlas datasets under covering and off-center prompt conditions, comparing against an automatic Attention U-Net, two Attention U-Net variants that we add ourselves to see how a conventional network behaves when it is also given the same box prompt (one that concatenates the prompt as an input channel, one that only looks at the prompt-cropped region), and SAM-Med2D; using Dice and IoU for region overlap, CBL for center-based localization, and HD95 for boundary distance, together with Monte Carlo cross-validation and computational efficiency measurements.
- Through a small-lesion analysis and an Attention U-Net-defined difficulty split, we show that the benefit of prompt-conditioned processing is concentrated exactly where automatic localization fails and where lesions are smallest.
- We describe a training-free self-assessment score, based on mask decisiveness and stability under prompt jitter, that ranks prediction reliability without ground truth and can also shortlist candidate regions for clinician review.

## II. RELATED WORK

For a long time, small-lesion segmentation in medical imaging has run into a simple fact: to segment a very small target well, a model must first narrow down where to look. In the fully automatic direction, this narrowing is learned implicitly inside the network. U-Net is the multi-scale encoder-decoder architecture that laid the foundation for biomedical segmentation [3], and Attention U-Net later added attention gates on the skip connections so the decoder could filter out irrelevant spatial responses more selectively [4]. Self-configuring pipelines such as nnU-Net go further still, automatically adapting preprocessing, network shape, and training schedule to each dataset [10]. These models raise automatic segmentation quality considerably, but still rest on a shared assumption: the network must localize and draw the mask from the input image alone, with no hint at all about where the lesion lies.

In parallel, interactive medical segmentation has also developed, starting from the observation that even a small amount of user guidance can substantially simplify the problem. Early methods used clicks, scribbles, or geodesic cues [5], [6]; later, bounding boxes were added as a fast, coarse form of prompt. More recently, interactive models such as ScribblePrompt can handle several prompt types within a single framework, across many biomedical datasets [11]. Broadly, these methods shift part of the localization burden from the model to the user. For very small X-ray lesions, that shift is especially valuable, since a clinician can usually point out the suspicious region coarsely even when an automatic system struggles to find it reliably in the whole image.

Promptable foundation models push this idea further. SAM pioneered turning prompt-driven segmentation into a general-purpose task [1]; SAM-Med2D and MedSAM brought it into medicine [2], [9]; adapter-based variants such as EMedSAM

**TABLE 1.** Representative prior work and the remaining gap addressed in this article.

| Work                                         | Principle                                              | Method                                                     | Performance Evidence                                                          | Pros                                                                                              | Cons                                                                                          |
|----------------------------------------------|--------------------------------------------------------|------------------------------------------------------------|-------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| U-Net [3]                                    | Fully convolutional encoder-decoder                    | Multi-scale skip-connected segmentation network            | Widely used as a canonical biomedical segmentation baseline                   | Simple, efficient, strong reference model                                                         | No prompt guidance; all localization must be automatic                                        |
| Attention U-Net [4]                          | Skip-feature filtering by attention gates              | U-Net with gated skip connections                          | Medical segmentation tasks reported with overlap metrics such as Dice and IoU | Stronger automatic baseline than plain U-Net                                                      | Attention is still learned without explicit prompt conditioning                               |
| mU-Net [10]                                  | Self-configuring automatic pipeline                    | Dataset-adaptive preprocessing, architecture, and training | Strong automatic results across many medical benchmarks                       | Robust automatic pipeline with little manual tuning                                               | No prompt mechanism; localization remains fully automatic                                     |
| DeepGloS and related interactive methods [6] | User guidance reduces search space                     | Geodesic or interaction-aware segmentation                 | Interactive medical segmentation under user input                             | Better reflects clinical interaction                                                              | Interaction mechanism can be task-specific or costly to maintain                              |
| ScribblePrompt [11]                          | General interactive biomedical segmentation            | One model for scribbles, clicks, and boxes across datasets | Multi-dataset interactive evaluation with several prompt types                | Flexible, realistic interaction types                                                             | General-purpose, not a lightweight radiograph-specific prompt prior                           |
| SAM-Med2D / MedSAM [2], [9]                  | Promptable foundation segmentation                     | Large pretrained model with prompt-driven mask decoding    | Multi-dataset medical evaluation with prompt-based metrics                    | Flexible and broadly applicable                                                                   | Heavy; prompt influence is not designed as a lightweight end-to-end radiograph-specific prior |
| EMedSAM / Med-SA [12], [13]                  | Adapter tuning of SAM for medical images               | Lightweight adapters on a frozen SAM backbone              | Medical segmentation with reduced adaptation cost                             | Cheaper specialization of a foundation model                                                      | Still a large SAM backbone; prompt acts mainly at the mask decoder                            |
| This work                                    | Prompt as a dense spatial prior throughout the network | Gaussian plateau prompt + PSG + CAD in a lightweight CNN   | BTXRD and FracAtlas; Dice, CBI, HD95, efficiency, and small-lesion analysis   | Lightweight, prompt-aware from encoder to decoder, evaluated under controlled prompt displacement | Still requires supervised masks and fixed-resolution preprocessing                            |

and Med-SA reduce the cost of specializing these models to medical data [12], [13]. These models are quite flexible, but they consume the prompt mainly at the mask decoder, carry a very large backbone, and none is tuned specifically for a single radiograph task. For very small bone lesions, this is exactly the gap: the prompt only marks where the mask goes, while no mechanism continues to shape feature extraction afterward, exactly where a weak signal is easily lost.

This article starts from that very gap. On grayscale bone radiographs, knowing roughly where the lesion is only solves half the problem; the harder part is holding on to faint evidence while still filtering out noise, after the search space has been narrowed. We argue that a lightweight prompt-guided CNN can do this better if the prompt is kept continuously active from the first encoder stage to the last decoder stage. In this article we study that setting directly, using SAM-Med2D as the promptable reference model; a broader comparison with newer promptable medical models is left to future work.

### III. METHOD

#### A. OFFLINE AND ONLINE FRAMEWORK

The system operates in two stages. In the offline stage, the model is trained on X-ray images that already have lesion polygons. Each polygon is turned into a training prompt by generating a simulated bounding box, then converting that box into a Gaussian-smoothed plateau heatmap. The network learns to map the (image, prompt) pair to a binary lesion mask. In the online stage, the user only needs to draw a coarse bounding box around the suspicious region on a new image; the same prompt-construction procedure is applied again, and the trained model produces the final predicted mask.

What sets this apart is where the prompt lives: instead of entering only once at the first layer, it is kept as a dense spatial prior, active throughout the encoder and decoder, which should help when the box is coarse, slightly off-center, or wrapped tightly around a very small lesion.

An interactive demonstration of the online stage is given in Figure 1: the user draws a box on the X-ray image, and the model returns the mask together with CAD's prompt-use score and the self-assessment score of Section III-F, both of which require no ground truth.

#### B. PROBLEM SETTING

Given a grayscale radiograph  $\mathbf{I} \in \mathbb{R}^{H \times W}$  and a user-provided bounding box  $\mathbf{B}$ , the goal is to predict a binary lesion mask  $\hat{\mathbf{M}}$ . Let  $\mathbf{H}(\mathbf{B})$  be the prompt heatmap derived from the box. PGA-UNet predicts

$$\hat{\mathbf{M}} = \mathbf{1} [f_{\text{PGA}}(\mathbf{I}, \mathbf{H}(\mathbf{B}); \theta) \geq 0.5]. \quad (1)$$

#### C. PROMPT REPRESENTATION


Instead of encoding the input prompt as a hard binary rectangle, we rasterize the interior of the box into a plateau map at the original image resolution, soften its edges with a fixed  $31 \times 31$  Gaussian kernel, and only then resize and pad this heatmap down to the target square resolution together with the image. The kernel size does not change with the target resolution ( $256 \times 256$  or  $512 \times 512$ ): applying it once, at the original resolution, before downsampling, keeps the effective blur consistent regardless of which square frame the network ultimately sees. The plateau covers exactly the area of the box (expanded if applicable), with no additional minimum margin added; coverage of the region of interest is already guaranteed by the box-expansion step, so no separate margin step is needed.

At evaluation time, the prompt is tested under two fixed conditions. The covering condition (`center_zoom`) scales the tight lesion box outward from its own center by a fixed factor of 3.0. The off-center condition (`center_shift`) applies the same scaling, then shifts the box by a fraction of the tight-box size (capped by `shift_ratio=0.5`), expanding it back if needed so the box still fully contains the lesion. At test time, this shift is fixed per lesion (seeded from the sample index), so every model is evaluated on exactly the same set of off-center boxes; during training, the shift is redrawn randomly at each iteration. For each dataset and resolution, we train a single model using `center_mixed`: `center_shift` is chosen with probability 0.8 and `center_zoom` with probability 0.2, while at test time the two conditions are scored separately. The task scope assumes the user has already identified the suspicious region and deliberately draws a box that covers the lesion; partial-coverage boxes, negative boxes, and off-target boxes all lie outside this controlled protocol, and whether a clinician can find the correct region in the first place is a separate question, outside the scope of this article.

Let  $(x_1, y_1, x_2, y_2)$  denote the tight lesion box, with  $w = x_2 - x_1$  and  $h = y_2 - y_1$ . The simulated covering box uses exactly the fixed center-scaling rule above, applied during both training and evaluation. The off-center box reuses that scaled box, then applies the deterministic test-time shift rule implemented in `dataset.py`, while still preserving lesion

### Interactive Demo: PGA-UNet (Prompt-Guided Attention U-Net)

Select a dataset and a test image with available ground truth, click two points to draw a lesion box, inspect the segmentation output, and compare all six metrics against the reference mask. Multiple prompt rounds can be added with Continue before switching to another image with Finish. Use Suggest prompts to have the model rank a batch of randomly sized/positioned boxes by the training free self-assessment score  $Q$ , then click a thumbnail to apply it like any other box.



**FIGURE 1.** Interactive demonstration of the online stage on a BTXRD case. The clinician draws a box on the left canvas; the right panel overlays the predicted mask, and the panel below reports the reference metrics together with the no-ground-truth CAD prompt-use score and the self-assessment score  $Q$ . A FracAtlas case is shown in the same layout in the supplementary material.

coverage. The scale factor, shift ratio, and Gaussian kernel size were all chosen from preliminary experiments on the training and validation data and then locked; the test split played no part whatsoever in choosing the protocol or any hyperparameter. These are simply fixed protocol values used throughout the article, not a globally optimal configuration. This simulation naturally cannot substitute for a prompt actually drawn by a radiologist, but it offers a controlled way to measure the model's sensitivity to a predefined prompt displacement.

In summary, prompt construction consists of three steps: (1) derive the lesion box from the annotated polygon or from user input, then expand it as described; (2) rasterize the expanded region into a plateau mask at the original image resolution; (3) smooth the plateau's edge with the fixed Gaussian kernel before resizing to the target resolution. This avoids the hard-edge dependence of a binary box while still keeping a clear spatial prior.

#### D. PROMPT SPATIAL GATE

Let  $\mathbf{x}^l$  denote the encoder feature map at level  $l$ . PSG modulates it with a prompt-derived attention map  $\mathbf{A}_{\text{PSG}}^l$ :

$$\tilde{\mathbf{x}}^l = \mathbf{x}^l \odot (1 + \alpha_{\text{PSG}}^l \mathbf{A}_{\text{PSG}}^l), \quad (2)$$

where  $\odot$  denotes element-wise multiplication, and  $\alpha_{\text{PSG}}^l$  is a learnable parameter, initialized at 0.1 and clamped to  $[0, 1]$ . The residual form preserves the original feature stream while emphasizing prompt-relevant regions.

#### E. CONDITIONAL ATTENTION DECODER


CAD extends decoder attention by conditioning the skip gate on prompt-encoded features. Let  $\mathbf{g}^l$  be the decoder gating feature and  $\mathbf{p}_{\text{enc}}^l$  the prompt encoding at the same scale. The conditioned gate is

$$\mathbf{g}'^l = \mathbf{g}^l + c^l \alpha_{\text{CAD}}^l w_l \mathbf{p}_{\text{enc}}^l, \quad (3)$$

where  $c^l$  is a prompt-use coefficient derived from the decoder,  $\alpha_{\text{CAD}}^l$  is learnable, and  $w_l$  is a fixed depth coefficient. As a result, the gate stays responsive to the prompt even when the box is not perfectly centered.

The prompt encoder inside CAD consists of two  $3 \times 3$  convolutional layers with instance normalization and ReLU. The prompt-use coefficient  $c^l$  is produced by global average pooling, followed by a  $1 \times 1$  convolution and a sigmoid; this coefficient is computed only inside the gate, used to scale the prompt term at that level, and is not exposed or analyzed as a separate model output in this article. The decoder's depth coefficients are fixed at  $\{1.0, 0.7, 0.4, 0.2\}$  from deep to shallow. With a feature scale of 4, the channel widths are  $\{16, 32, 64, 128, 256\}$  from the shallowest encoder stage to the bottleneck, and the whole model has about 2.95 million parameters. These protocol values were all chosen from preliminary experiments on the training and validation sets, held fixed throughout the main evaluation, with the test split never involved in tuning; we make no claim of optimality. The CAD mixing coefficient is parameterized through a sigmoid and initialized near 0.3.





**FIGURE 2.** Overview of PGA-UNet. Prompt information derived from the bounding box is injected into encoder features through PSG and reused in decoder skip attention through CAD.

In summary, the method consists of four steps: (1) pre-process the image and prompt to the same square resolution; (2) encode the prompt-relevant spatial prior with PSG in the encoder; (3) propagate prompt-conditioned context through CAD in the decoder; (4) predict the final binary lesion mask.

#### F. TRAINING-FREE SELF-ASSESSMENT SCORE

Besides the predicted mask, PGA-UNet also returns a scalar  $Q$ , estimating how trustworthy a prediction is without needing the true answer.  $Q$  is not a learned quantity; it is computed only from the model's own output at inference time, as the average of two cues.

The first cue,  $S_{\text{sharp}}$ , measures how decisive the predicted mask is. Let  $p$  be the sigmoid probability map and  $\Omega = \{p > 0.5\}$  the predicted positive region. Then

$$S_{\text{sharp}} = \text{clip}_{[0,1]} \left( \frac{1}{|\Omega|} \sum_{u \in \Omega} (2p(u) - 1) \right). \quad (4)$$

A mask with high confidence, i.e., probabilities close to 1, scores close to 1; conversely, a mask with many pixels sitting near the decision threshold scores lower.

The second cue,  $S_{\text{stab}}$ , measures how stable the prediction is under mild prompt perturbation. The box is randomly nudged  $K$  times, each time with a small shift and scale change; the model is rerun for each nudged box, and  $S_{\text{stab}}$  is exactly the average IoU between the original binary mask  $\hat{M}$  and the masks obtained from those nudges,  $\hat{M}_k$ :

$$S_{\text{stab}} = \frac{1}{K} \sum_{k=1}^K \text{IoU}(\hat{M}, \hat{M}_k). \quad (5)$$

A prediction that stays nearly unchanged when the box is nudged slightly is considered more trustworthy than one that collapses or drifts elsewhere.

The final score is  $Q = \frac{1}{2}(S_{\text{sharp}} + S_{\text{stab}})$ , with  $K = 6$ , a maximum shift of 0.12 of the box side, and a scale factor within  $\pm 0.10$ . It should be emphasized that  $Q$  is only a heuristic ranking signal, not a calibrated probability: a value of 0.7 does not mean a 70% probability that the mask is correct.

**TABLE 2.** Image-level train / validation / test splits. Each lesion polygon is one promptable sample.

| Dataset   | Images |     |      | Lesion polygons |     |      |
|-----------|--------|-----|------|-----------------|-----|------|
|           | Train  | Val | Test | Train           | Val | Test |
| BTXRD     | 1,493  | 187 | 187  | 1,848           | 238 | 232  |
| FracAtlas | 573    | 72  | 72   | 730             | 95  | 92   |

We reuse this same score to shortlist candidate regions when no prompt is given: a set of boxes with varying size and position is sampled, each box is scored with  $Q$ , and the highest-scoring few are returned for the clinician to accept, adjust, or discard. This is only a review-support tool, not an automatic detector.

## IV. EXPERIMENTAL SETUP

### A. DATASETS

We run our experiments on two public bone X-ray datasets. BTXRD is a primary bone tumor dataset with polygon annotations supporting classification, localization, and segmentation [7]. FracAtlas is a fracture dataset with the same label types [8]; we clean and renormalize the polygons before use. Both datasets are split at the image level, so that every polygon belonging to the same image always falls in the same split (Table 2).

Because the public metadata do not provide enough information to verify strict patient-level separation, images from the same patient could in principle still fall into different splits; we regard this as a limitation of this evaluation. In image-level evaluation, if an image has multiple lesions, the model is run separately for each box, and the resulting probability maps are merged by a pixelwise maximum before thresholding at 0.5; the ground-truth masks are likewise merged by union.

Input images are resized isotropically to preserve aspect ratio, then padded to a square resolution of  $256 \times 256$  or  $512 \times 512$ . Images, masks, and prompt maps all undergo this same geometric preprocessing.

### B. TRAINING DETAILS

All CNN models were implemented in PyTorch and trained on a single NVIDIA Tesla T4 GPU with 16 GB memory. PGA-UNet and Attention U-Net were trained separately for each dataset and each resolution; in other words, this article studies one shared architecture family instantiated as separate models, with separate parameter sets for BTXRD and FracAtlas, rather than training a single joint model on both. We used the AdamW optimizer with learning rate  $10^{-4}$  and weight decay  $10^{-4}$ ; the loss was the sum of binary cross-entropy and Dice loss. Because the segmentation targets are often small, we additionally tested training PGA-UNet at  $512 \times 512$  with the Dice term replaced by a size-conditioned Tversky loss, and in a separate run, by a Focal Dice loss, with every other setting held fixed. These two replacements are mutually

exclusive and are reported later only as a negative control on the training objective. Batch size was 4, training ran for at most 150 epochs, with early stopping at patience 15. All main and ablation experiments share the same fixed training seed, 22120196, which seeds Python, NumPy, and PyTorch on both CPU and CUDA. The primary model-selection criterion was image-level merged Dice, measured on the validation set under the `center_shift` condition.

During training, we applied augmentation consisting of horizontal flipping and random rotation within  $\pm 15^\circ$ . PGA-UNet is trained with `center_mixed: center_shift` is chosen with probability 0.8 and `center_zoom` with probability 0.2. Because the prompts are all simulated from labeled lesion regions, this is a prompted-segmentation protocol under controlled conditions; the article should be read accordingly, namely as a study of simulated lesion-covering boxes around user-identified lesions, not as an evaluation of unconstrained lesion detection or of boxes fully drawn by a clinician in real practice.

For the fine-tuned SAM-Med2D comparison, the pre-trained image encoder was frozen, with only its lightweight Adapter layers trained further, while the prompt encoder and mask decoder were fully updated on each dataset, under the same split protocol. SAM-Med2D was fine-tuned with box prompts only, using the same center-scaled covering and off-center protocol as PGA-UNet; the point-prompt branch and iterative point refinement were both disabled. This box construction differs from the small random-jitter `get_boxes_from_mask` sampling used by the original SAM-Med2D authors, which we did not apply here. Both validation during SAM-Med2D fine-tuning and the final test-split evaluation report both prompt conditions separately; batch size was 4, at most 150 epochs, early stopping patience 30, and learning rate  $10^{-5}$ . This matched protocol reduces differences caused by prompt distribution and resolution, but cannot eliminate differences in model scale, initialization, or optimization.

We additionally ran stability experiments using Monte Carlo cross-validation, that is, repeated random image-level splits, as opposed to strict non-overlapping  $k$ -fold partitioning. These experiments keep the training seed fixed at 22120196 and vary only the split seed, across the values 1, 2, 3, and 4.

Two details are worth noting when interpreting the evaluation results. First, prompt construction itself does not depend on resolution: the plateau heatmap is built and smoothed with a fixed  $31 \times 31$  Gaussian kernel at the original image resolution, and only then resized and padded down to  $256 \times 256$  or  $512 \times 512$  together with the image. In other words, the same absolute blur is compressed differently at the two target resolutions, rather than being parameterized separately for each. Second, the Monte Carlo experiments are only auxiliary stability evidence, not intended to replace the fixed train/validation/test split used throughout the main results tables.

### C. BASELINES AND METRICS

We compare PGA-UNet against two main baselines: Attention U-Net [4], serving as the automatic baseline, and SAM-Med2D [2], serving as the prompt-based baseline. Attention U-Net is trained without a prompt, to show how hard the task becomes when localization must be fully automatic. SAM-Med2D is evaluated at  $256 \times 256$ , in fine-tuned form, using the same dataset splits and the same box prompts as PGA-UNet, making it the article's primary prompt-matched comparison.

To isolate whether PGA-UNet's advantage comes from its architecture or simply from having access to a box prompt at all, we add two conventional prompt-matched baselines of our own, both built on the same Attention U-Net backbone and given exactly the same box prompt that PGA-UNet receives, but without PGA-UNet's Gaussian-prior heatmap, Prompt Spatial Gate, or Conditional Attention Decoder. The first baseline, Attention U-Net with a prompt channel, concatenates the prompt heatmap as a second input channel and predicts on the full image through a plain skip-connection decoder. The second baseline, Attention U-Net on prompt crops, narrows the network's field of view: the input image is cropped to the prompt box before resizing to the model's input resolution, the plain Attention U-Net predicts a mask on that crop alone, and the result is pasted back into the full frame for evaluation. Both baselines are trained under the same covering and off-center box protocol as PGA-UNet, so any remaining gap reflects a difference in architecture rather than an unfair prompt. When comparing models against each other, we use the pasted-back full-frame metric for the prompt-crop baseline, not the diagnostic value measured inside the crop frame.

Every comparison table in this article reports Dice and IoU for region overlap, CBL for localization, and HD95 for boundary distance. Dice and IoU are fairly strongly correlated overlap measures; we place them side by side only for the reader's convenience in cross-checking. On the small-lesion subsets specifically, IoU has a low absolute value and is quite sensitive to a few boundary pixels, so Dice is the primary overlap metric used there. Precision and recall are kept in the released logs that accompany the article. Let  $(x_p, y_p)$  be the predicted mask centroid,  $(x_g, y_g)$  the ground-truth centroid, and  $D_{GT}$  the diagonal length of the ground-truth bounding box; CBL is defined as follows:

$$\text{CBL} = \max \left( 0, 1 - \frac{\sqrt{(x_p - x_g)^2 + (y_p - y_g)^2}}{D_{GT}} \right). \quad (6)$$

HD95 is computed from bidirectional boundary distances. When both masks are empty, HD95 is set to 0; if only one mask is empty, HD95 is set to the input side length  $S$  (256 or 512 pixels). No sample is discarded from the computation. For comparisons spanning resolutions, we additionally report HD95 in normalized form,  $\text{HD95}/(\sqrt{2}S)$ , since a raw pixel distance cannot be directly compared across the 128, 256, and 512 frames. Some subgroup comparisons mix models trained at  $256 \times 256$  and  $512 \times 512$ ; in that case, everything

is scored together in a shared  $512 \times 512$  frame, with the  $256 \times 256$  predictions upsampled before scoring, so every model is measured against exactly the same ground-truth mask, rather than in its own native frame.

For the training-free self-assessment score, each lesion prompt yields a score  $Q$  together with the true thresholded Dice of its own predicted mask; we report the Pearson and Spearman correlation between the two, separately for the covering and off-center conditions. For the candidate-region shortlist, each test image is sampled with 50 boxes whose side length is uniform in  $[0.15, 0.55]$  of the frame, each box is scored by  $Q$ , and the five highest-scoring boxes are kept; each retained box is then scored by coverage, the fraction of ground-truth lesion pixels it contains, and no predicted mask enters this coverage measure.

#### D. SCOPE OF ANALYSIS

This article focuses only on prompt-guided segmentation itself. A lesion-screening classification stage was explored in a separate study and is not part of this article. The self-assessment score and the candidate-region shortlist are treated only as auxiliary analyses, within the same narrow scope already stated in Section III-F, and are not extended any further here. To keep the article compact, we retain only the subgroup analyses that most directly support the main claim: the role of prompt-guided localization, sensitivity to controlled prompt displacement as defined above, and effectiveness on very small lesions. For each dataset, the small-lesion subgroup is defined as the 50 test images with the smallest total lesion area. In the three-resolution small-lesion comparison, this same set of images is held fixed and evaluated in turn at 128, 256, and 512; predictions and ground truth are merged at image level at each resolution, HD95 is reported both in native pixels and in the normalized form, and the degradation is measured relative to the 512 result. It should also be reiterated: partial-coverage boxes, negative boxes, and off-target boxes are outside the scope of this study, so the article's claims do not rely on them.

#### V. RESULTS

Unless stated otherwise, every table reports current-protocol numbers: center\_mixed training with 80% center\_shift and 20% center\_zoom, 150 epochs, model selection by image-level merged validation Dice under center\_shift, and image-level merged test evaluation. The two test prompt conditions, covering (center\_zoom) and off-center (center\_shift), are reported separately whenever a table carries both.

##### A. MAIN RESULTS ACROSS DATASETS AND RESOLUTIONS

PGA-UNet was trained from scratch at  $128 \times 128$ ,  $256 \times 256$ , and  $512 \times 512$  on each dataset, giving six models. Table 3 reports them under both prompt conditions.

Higher resolution helps at every step: Dice consistently follows  $512 > 256 > 128$  on both datasets and under both prompt conditions, with FracAtlas gaining more, which

**TABLE 3. PGA-UNet across resolutions, image-level merged. Dice, IoU, CBL and HD95 (native px and normalized HD95/( $\sqrt{2.5}$ )) given as covering / off-center.**

| Dataset   | Res. | Dice                 | IoU                  | HD95 <sub>px</sub> | HD95 <sub>n</sub>    | CBL                  |
|-----------|------|----------------------|----------------------|--------------------|----------------------|----------------------|
| BTXRD     | 128  | 0.714 / 0.717        | 0.572 / 0.577        | 5.8 / 5.3          | 0.032 / 0.029        | 0.859 / 0.857        |
|           | 256  | 0.762 / 0.743        | 0.633 / 0.611        | 10.9 / 12.8        | 0.030 / 0.035        | 0.902 / 0.885        |
|           | 512  | <b>0.782 / 0.774</b> | <b>0.661 / 0.652</b> | 22.8 / 24.5        | 0.031 / 0.034        | <b>0.918 / 0.912</b> |
| FracAtlas | 128  | 0.596 / 0.542        | 0.445 / 0.389        | 8.1 / 9.5          | 0.044 / 0.053        | 0.849 / 0.790        |
|           | 256  | 0.629 / 0.594        | 0.470 / 0.436        | 10.6 / 11.7        | 0.029 / 0.032        | 0.895 / 0.849        |
|           | 512  | <b>0.727 / 0.713</b> | <b>0.585 / 0.570</b> | 10.5 / 10.4        | <b>0.015 / 0.014</b> | <b>0.913 / 0.897</b> |

**TABLE 4. Comparison at  $512 \times 512$  under covering prompts. Attention U-Net (no prompt) receives no box; the other rows use the same box PGA-UNet receives. 'Dice is each row minus PGA-UNet.**

| Dataset   | Model                            | Dice         | $\Delta$ Dice | IoU          | HD95        | CBL          |
|-----------|----------------------------------|--------------|---------------|--------------|-------------|--------------|
| BTXRD     | Attention U-Net (no prompt)      | 0.517        | -0.265        | 0.418        | 120.4       | 0.620        |
|           | Attention U-Net + prompt channel | 0.760        | -0.022        | 0.635        | 33.0        | 0.895        |
|           | Attention U-Net + prompt crop    | 0.741        | -0.041        | 0.615        | 24.7        | 0.912        |
|           | PGA-UNet                         | <b>0.782</b> | —             | <b>0.661</b> | <b>22.8</b> | <b>0.918</b> |
| FracAtlas | Attention U-Net (no prompt)      | 0.342        | -0.385        | 0.253        | 139.9       | 0.478        |
|           | Attention U-Net + prompt channel | 0.593        | -0.134        | 0.434        | 22.9        | 0.887        |
|           | Attention U-Net + prompt crop    | 0.590        | -0.137        | 0.451        | 22.0        | 0.865        |
|           | PGA-UNet                         | <b>0.727</b> | —             | <b>0.585</b> | <b>10.5</b> | <b>0.913</b> |

is reasonable since its inherently thin fracture lesions lose proportionally more detail as the image shrinks. Because HD95 in raw pixels cannot be compared across differently sized frames, we also report the normalized form: it stays near 0.03 across the BTXRD resolutions and decreases on FracAtlas as the frame grows, showing that boundary error does not worsen at higher resolution. Note that each model is evaluated only on the dataset it was trained on; the article makes no claim about cross-dataset transfer.

##### B. COMPARISON AGAINST AUTOMATIC AND PROMPT-MATCHED ATTENTION U-NET

With no support at all in finding the lesion, Attention U-Net reaches only 0.517 Dice on BTXRD and 0.342 on FracAtlas at  $512 \times 512$ , with HD95 exceeding 120 pixels on both datasets (Table 4). This number should not be read as a verdict on the architecture, since the model itself has never seen the box; it only shows how hard automatic localization is, exactly the part of the job that a prompt helps remove.

Given the same box, the plain Attention U-Net does considerably better, but a gap to PGA-UNet remains (Table 4,  $\Delta$ Dice column). PGA-UNet surpasses both prompt-matched variants on both datasets, with a small margin on BTXRD but a much wider one on FracAtlas, and lower HD95 than either.


So localization accounts for most of the difficulty here, but the box alone does not fully explain the results. The way PGA-UNet exploits that box still produces an additional, consistent margin over a conventional network given the same input.

##### C. ATTENTION U-NET TOP-DICE AND BOTTOM-DICE SUBSETS

To see where that margin comes from, we ranked the test images by Attention U-Net Dice, kept the top 50 and bottom 50 per dataset, and re-ran all four models on those exact images (Table 5, covering prompts). On the top-50, the

**TABLE 5.** Top-Dice and bottom-Dice subsets at  $512 \times 512$ , 50 images per group, covering prompts. Groups are defined by plain Attention U-Net.


| Dataset   | Group     | Model             | Dice         | IoU          | HD95  | CBL          |
|-----------|-----------|-------------------|--------------|--------------|-------|--------------|
| BTXRD     | Top-50    | Attention U-Net   | 0.872        | 0.775        | 35.0  | 0.945        |
|           |           | AttUNet + channel | 0.846        | 0.740        | 24.3  | 0.942        |
|           |           | AttUNet + crop    | 0.844        | 0.739        | 27.5  | 0.946        |
|           |           | PGA-UNet          | 0.862        | 0.763        | 19.2  | 0.944        |
|           | Bottom-50 | Attention U-Net   | 0.031        | 0.016        | 277.8 | 0.105        |
|           |           | AttUNet + channel | 0.668        | 0.541        | 43.3  | 0.831        |
|           |           | AttUNet + crop    | 0.652        | 0.518        | 17.6  | 0.875        |
|           |           | PGA-UNet          | <b>0.712</b> | <b>0.589</b> | 23.2  | <b>0.879</b> |
| FracAtlas | Top-50    | Attention U-Net   | 0.493        | 0.365        | 85.7  | 0.640        |
|           |           | AttUNet + channel | 0.605        | 0.447        | 26.6  | 0.887        |
|           |           | AttUNet + crop    | 0.658        | 0.515        | 24.9  | 0.881        |
|           |           | PGA-UNet          | <b>0.754</b> | <b>0.615</b> | 10.6  | <b>0.927</b> |
|           | Bottom-50 | Attention U-Net   | 0.169        | 0.106        | 182.3 | 0.315        |
|           |           | AttUNet + channel | 0.574        | 0.416        | 25.3  | 0.875        |
|           |           | AttUNet + crop    | 0.509        | 0.369        | 26.0  | 0.834        |
|           |           | PGA-UNet          | <b>0.687</b> | <b>0.535</b> | 11.8  | <b>0.896</b> |

**FIGURE 3.** Attention U-Net (no prompt) vs PGA-UNet on BTXRD, Attention U-Net top-Dice case (IMG001375). Both models do well on an easy case.

prompt-guided models perform quite similarly. In contrast, on the bottom-50, Attention U-Net collapses outright, down to 0.031 on BTXRD and 0.169 on FracAtlas, while PGA-UNet still holds 0.712 and 0.687. The two conventional prompt baselines also recover somewhat, but less: on the bottom group PGA-UNet still leads them by 0.04 to 0.07 Dice on BTXRD and by 0.10 to 0.19 on FracAtlas. This shows that having a prompt is necessary but not sufficient; the benefit of prompt-conditioned processing is concentrated exactly where automatic localization fails. It should also be noted that these two groups are defined entirely by Attention U-Net, not by any other objective difficulty criterion.

#### D. MATCHED COMPARISON AGAINST SAM-MED2D

Both models receive the same box at  $256 \times 256$  (Table 6). In zero-shot mode, SAM-Med2D barely works on bone X-ray images: only 0.255 Dice on BTXRD and 0.262 on FracAtlas. After fine-tuning on the same split and the same prompt protocol, this rises to 0.630 and 0.563; however, PGA-UNet at the same resolution reaches 0.762 and 0.629 under covering prompts, ahead by 0.132 and 0.066 Dice, with a similar gap in CBL. SAM-Med2D pulls ahead in exactly one place: boundary distance on FracAtlas, where its fine-tuned HD95 of 8.0 pixels beats PGA-UNet's 10.6. The change from covering to

**FIGURE 4.** Attention U-Net (no prompt) vs PGA-UNet on BTXRD, Attention U-Net bottom-Dice case (IMG000791). Without a prompt, Attention U-Net's prediction is badly mislocated; the prompt-conditioned PGA-UNet still recovers the lesion.**TABLE 6.** Comparison with SAM-Med2D at  $256 \times 256$ , image-level merged, covering prompts.

| Dataset   | Model                | Dice         | IoU          | HD95        | CBL          |
|-----------|----------------------|--------------|--------------|-------------|--------------|
| BTXRD     | SAM-Med2D zero-shot  | 0.255        | 0.156        | 41.2        | 0.734        |
|           | SAM-Med2D fine-tuned | 0.630        | 0.497        | 17.3        | 0.836        |
|           | PGA-UNet             | <b>0.762</b> | <b>0.633</b> | <b>10.9</b> | <b>0.902</b> |
| FracAtlas | SAM-Med2D zero-shot  | 0.262        | 0.156        | 21.3        | 0.752        |
|           | SAM-Med2D fine-tuned | 0.563        | 0.417        | <b>8.0</b>  | 0.846        |
|           | PGA-UNet             | <b>0.629</b> | <b>0.470</b> | 10.6        | <b>0.895</b> |

off-center prompts is discussed separately in Section V-F.

#### E. PERFORMANCE ON SMALL LESIONS

This is exactly where keeping the prompt's features active throughout the network should matter most: the small-lesion subset consists of the 50 test images per dataset with the smallest total lesion area, ranging from 7 to 433 pixels on BTXRD and 145 to 944 pixels on FracAtlas. All scores here are computed in a shared  $512 \times 512$  frame.

##### a: Against SAM-Med2D at the matched resolution

SAM-Med2D struggles clearly at this lesion size (Table 7). Even after fine-tuning at 256, it reaches only 0.256 Dice on BTXRD and 0.371 on FracAtlas; in zero-shot mode it is close to nothing. PGA-UNet at the same 256 resolution still holds 0.705 and 0.648, with far lower HD95.

##### b: Against prompt-matched Attention U-Net at 512

Given the same box at 512, the two conventional prompt baselines trail PGA-UNet by about 0.04 on BTXRD and by 0.11 to 0.16 on FracAtlas, while prompt-free Attention U-Net drops to only about 0.28 (Table 8). The prompt-crop variant reaches a lower HD95 on BTXRD, perhaps because its tight crop suits a compact target well, but this does not carry over to FracAtlas.


##### c: Resolution on the small-lesion stems

It is exactly on this group of images that resolution shows its clearest effect. On the fixed small-lesion stems, PGA-UNet's



**TABLE 7.** Small-lesion subset, PGA-UNet vs SAM-Med2D at the matched  $256 \times 256$  resolution. 50 images per dataset, covering prompts, scored in a shared  $512 \times 512$  frame.

| Dataset   | Model                    | Dice         | IoU          | HD95        | CBL          |
|-----------|--------------------------|--------------|--------------|-------------|--------------|
| BTXRD     | SAM-Med2D zero-shot 256  | 0.139        | 0.082        | 62.4        | 0.515        |
|           | SAM-Med2D fine-tuned 256 | 0.256        | 0.170        | 123.8       | 0.561        |
|           | PGA-UNet 256             | <b>0.705</b> | <b>0.559</b> | <b>14.7</b> | <b>0.877</b> |
| FracAtlas | SAM-Med2D zero-shot 256  | 0.205        | 0.123        | 33.5        | 0.659        |
|           | SAM-Med2D fine-tuned 256 | 0.371        | 0.256        | 62.1        | 0.736        |
|           | PGA-UNet 256             | <b>0.648</b> | <b>0.491</b> | <b>20.6</b> | <b>0.897</b> |

**FIGURE 5.** Small-lesion case on BTXRD (IMG000932): SAM-Med2D zero-shot, SAM-Med2D fine-tuned, and PGA-UNet, all at  $256 \times 256$ . Both SAM-Med2D variants miss the lesion entirely; PGA-UNet does not.

Dice falls from 0.766 at 512 to 0.631 at 128 on BTXRD, and from 0.720 to 0.609 on FracAtlas; the loss from 512 to 128 is 0.135 and 0.111 respectively, markedly larger than the loss on the full test set (Table 9).

Overall, on precisely these smallest lesions, both prompt conditioning and input resolution matter substantially more than they do on the full test set.

### F. SENSITIVITY TO PROMPT DISPLACEMENT


Table 10 shows what happens when the box is shifted off-center under one fixed rule. PGA-UNet loses only 0.008 Dice on BTXRD and 0.014 on FracAtlas, with CBL dropping by 0.006 and 0.015 respectively; on the small-lesion subset the loss stays under 0.01 for both datasets. The conventional prompt baselines lose a comparable amount or slightly more, and fine-tuned SAM-Med2D loses about 0.03. This is only sensitivity to one displacement rule fixed in advance. We have not yet tested partial-coverage boxes or off-target boxes, so this number says nothing about boxes a clinician might freely draw in practice.

### G. ABLATION STUDY

The ablation drops one factor at a time from the full model (Gaussian prompt combined with PSG and CAD), and sepa-

**TABLE 8.** Small-lesion subset at  $512 \times 512$ , PGA-UNet vs Attention U-Net baselines. 50 images per dataset, covering prompts.

| Dataset   | Model                            | Dice         | IoU          | HD95       | CBL          |
|-----------|----------------------------------|--------------|--------------|------------|--------------|
| BTXRD     | Attention U-Net (no prompt)      | 0.283        | 0.219        | 178.2      | 0.283        |
|           | Attention U-Net + prompt channel | 0.728        | 0.591        | 14.2       | 0.879        |
|           | Attention U-Net + prompt crop    | 0.724        | 0.602        | <b>5.5</b> | <b>0.909</b> |
|           | PGA-UNet                         | <b>0.766</b> | <b>0.637</b> | 14.0       | 0.897        |
| FracAtlas | Attention U-Net (no prompt)      | 0.291        | 0.216        | 139.4      | 0.401        |
|           | Attention U-Net + prompt channel | 0.610        | 0.449        | 21.9       | 0.888        |
|           | Attention U-Net + prompt crop    | 0.564        | 0.425        | 22.6       | 0.855        |
|           | PGA-UNet                         | <b>0.720</b> | <b>0.578</b> | <b>8.3</b> | <b>0.907</b> |

**FIGURE 6.** Small-lesion case on BTXRD (IMG000932): Attention U-Net (no prompt), Attention U-Net with prompt channel, Attention U-Net with prompt crop, and PGA-UNet, all at  $512 \times 512$ .


ately tests replacing the Gaussian prompt with a binary box (Table 11). The full model achieves the best Dice in three of the four dataset-and-prompt-condition combinations. The only exception is BTXRD under covering prompts, where the binary-box variant is 0.006 Dice higher, a margin within the noise of a single split; in that same configuration, the full model still keeps the lowest HD95 and the highest CBL. The Gaussian prompt's advantage is clearest under off-center prompts, where the full model wins on both datasets, and especially on FracAtlas, where it wins every configuration on every metric. The goal of this experiment is to test the combination as a system, not to claim that any single internal mechanism is the most important. These are single-split results.

### H. MONTE CARLO CROSS-VALIDATION

Because a fixed split could by chance flatter or hurt the method, we retrained PGA-UNet at  $512 \times 512$  on four differ-

**TABLE 9.** PGA-UNet on the small-lesion subset across resolutions, 50 fixed stems per dataset. Dice, CBL and HD95 (native px / normalized) as covering / off-center. ' is the covering Dice loss relative to 512.

| Dataset   | Res. | Dice                 | IoU                  | CBL                  | HD95 <sub>px / n</sub> | $\Delta$ Dice |
|-----------|------|----------------------|----------------------|----------------------|------------------------|---------------|
| BTXRD     | 128  | 0.631 / 0.662        | 0.488 / 0.517        | 0.738 / 0.758        | 3.9 / 0.022            | −0.135        |
|           | 256  | 0.714 / 0.695        | 0.570 / 0.548        | 0.869 / 0.849        | 7.2 / 0.020            | −0.052        |
|           | 512  | <b>0.766 / 0.757</b> | <b>0.637 / 0.627</b> | <b>0.897 / 0.890</b> | 14.0 / 0.019           | —             |
| FracAtlas | 128  | 0.609 / 0.543        | 0.462 / 0.393        | 0.828 / 0.757        | 8.9 / 0.049            | −0.111        |
|           | 256  | 0.642 / 0.610        | 0.484 / 0.451        | 0.887 / 0.845        | 10.3 / 0.029           | −0.079        |
|           | 512  | <b>0.720 / 0.712</b> | <b>0.578 / 0.570</b> | <b>0.907 / 0.893</b> | 8.3 / 0.012            | —             |



**FIGURE 7.** PGA-UNet on the same small-lesion case on BTXRD (IMG000932) at 128 × 128, 256 × 256, and 512 × 512. The prediction tightens around the lesion as resolution increases.

ent random image-level splits, keeping the training seed fixed at 22120196 and changing only the split seed (Table 12). The results show that Dice varies by at most 0.011 across splits, for each dataset and each prompt condition; CBL varies by at most 0.013. Notably, the repeated-split mean on FracAtlas is 0.733, slightly higher than the 0.727 obtained on the fixed split used in the article, suggesting that if anything, the split chosen for reporting is somewhat harder than typical. This is evidence that PGA-UNet is stable across different data splits, not evidence that it outperforms any baseline.

## I. EFFICIENCY


PGA-UNet is very small, with all 2.955 million parameters trainable (Table 13). At the matched 256 × 256 resolution, it is about 92 times lighter than fine-tuned SAM-Med2D in parameter count, 12 times lighter in FLOPs, 215 times smaller in storage size, and 6 times faster on GPU. These are only cost figures measured on one specific implementation, one specific GPU, and the two resolutions used, and should not be generalized further.

## J. EXPLORATORY COMPARISON OF SEGMENTATION LOSSES

Because the segmentation targets in this article are small, we tried replacing the Dice term with a size-conditioned Tversky

**TABLE 10.** Sensitivity to prompt displacement. Every model under covering (z) and off-center (s) prompts. Attention U-Net (no prompt) is prompt-invariant. AttUNet rows and PGA-UNet at 512 × 512; SAM-Med2D fine-tuned at 256 × 256, image-level merged.

| Dataset   | Model                            | Dice z/s             | $\Delta$ Dice | IoU z/s              | HD95 z/s           | CBL z/s              |
|-----------|----------------------------------|----------------------|---------------|----------------------|--------------------|----------------------|
| BTXRD     | Attention U-Net (no prompt)      | 0.517 / 0.517        | 0.000         | 0.418 / 0.418        | 120.4 / 120.4      | 0.620 / 0.620        |
|           | Attention U-Net + prompt channel | 0.760 / 0.748        | −0.012        | 0.635 / 0.626        | 33.0 / 32.6        | 0.895 / 0.890        |
|           | SAM-Med2D fine-tuned             | 0.741 / 0.733        | −0.008        | 0.615 / 0.607        | 24.7 / 26.4        | 0.912 / 0.896        |
|           | PGA-UNet                         | 0.630 / 0.601        | −0.029        | 0.497 / 0.471        | 17.3 / 17.5        | 0.836 / 0.813        |
|           | PGA-UNet                         | <b>0.782 / 0.774</b> | −0.008        | <b>0.661 / 0.652</b> | 22.8 / 24.5        | <b>0.918 / 0.912</b> |
| FracAtlas | Attention U-Net (no prompt)      | 0.342 / 0.342        | 0.000         | 0.253 / 0.253        | 139.9 / 139.9      | 0.478 / 0.478        |
|           | Attention U-Net + prompt channel | 0.593 / 0.588        | −0.005        | 0.434 / 0.428        | 22.9 / 24.0        | 0.887 / 0.860        |
|           | Attention U-Net + prompt crop    | 0.590 / 0.563        | −0.027        | 0.451 / 0.428        | 22.0 / 23.8        | 0.865 / 0.833        |
|           | SAM-Med2D fine-tuned             | 0.563 / 0.532        | −0.031        | 0.417 / 0.389        | 8.0 / 8.4          | 0.846 / 0.816        |
|           | PGA-UNet                         | <b>0.727 / 0.713</b> | −0.014        | <b>0.585 / 0.570</b> | <b>10.5 / 10.4</b> | <b>0.913 / 0.897</b> |



**FIGURE 8.** PGA-UNet on the same BTXRD case (IMG000020) under a covering prompt (top) and an off-center prompt (bottom); Dice is unchanged at 0.942.

loss, and in another trial, with a Focal Dice loss, with every other setting held fixed (Table 14). Neither option worked better: the default Dice + BCE is always best or tied on both datasets and both prompt conditions, while Focal Dice is even a clear step backward on FracAtlas, with HD95 exceeding 100 pixels. We treat this as a negative control: under the current protocol, the default objective remains the better choice, and neither of the two alternatives is adopted.

## K. FAILURE MODES


PGA-UNet still makes mistakes, and these mistakes are not random but recur in a few patterns. In qualitative review, low Dice tends to cluster on lesions that are elongated or branching, lesions located in cluttered anatomy such as the shoulder or clavicle, or lesions that barely stand out from the surrounding bone. In such cases, even a good covering box does not help if the local image evidence itself is too weak.

## L. TRAINING-FREE SELF-ASSESSMENT AND CANDIDATE SUGGESTION

Even without any ground truth, the score  $Q$  presented in Section III-F still tracks the true Dice of each prompt fairly closely (Table 15). The Spearman correlation between  $Q$  and the true Dice reaches 0.70 (covering) and 0.60 (off-center) on BTXRD, and 0.65 and 0.48 on FracAtlas. Both component cues contribute meaningfully: mask decisiveness alone already reaches a Spearman of nearly 0.4, jitter stability reaches about 0.5 to 0.6, and averaging the two cues is always at least as good as either alone. As a control, we also tried training a

**TABLE 11.** Ablation at  $512 \times 512$ , image-level merged, single split. Dice, IoU, HD95 and CBL as covering / off-center.

| Dataset   | Configuration                        | Dice                 | IoU                  | HD95               | CBL                  |
|-----------|--------------------------------------|----------------------|----------------------|--------------------|----------------------|
| BTXRD     | CAD only                             | 0.771 / 0.759        | 0.649 / 0.635        | 23.1 / 24.2        | 0.915 / 0.905        |
|           | PSG only                             | 0.779 / 0.771        | 0.660 / 0.651        | 29.4 / 28.4        | 0.908 / 0.902        |
|           | PSG + vanilla attention gate         | 0.769 / 0.757        | 0.645 / 0.633        | 26.4 / 28.7        | 0.906 / 0.896        |
|           | Full model + binary box prompt       | <b>0.788 / 0.770</b> | <b>0.668 / 0.651</b> | 23.1 / 28.5        | 0.914 / 0.900        |
|           | Full PGA-UNet (Gaussian + PSG + CAD) | 0.782 / <b>0.774</b> | 0.661 / <b>0.652</b> | <b>22.8 / 24.5</b> | <b>0.918 / 0.912</b> |
| FracAtlas | CAD only                             | 0.696 / 0.685        | 0.551 / 0.540        | 18.0 / <b>12.0</b> | 0.903 / 0.885        |
|           | PSG only                             | 0.635 / 0.611        | 0.479 / 0.459        | 14.5 / 14.8        | 0.894 / 0.857        |
|           | PSG + vanilla attention gate         | 0.690 / 0.691        | 0.548 / 0.547        | 14.0 / 14.1        | 0.892 / 0.881        |
|           | Full model + binary box prompt       | 0.668 / 0.657        | 0.514 / 0.505        | 12.0 / 13.2        | 0.904 / 0.886        |
|           | Full PGA-UNet (Gaussian + PSG + CAD) | <b>0.727 / 0.713</b> | <b>0.585 / 0.570</b> | <b>10.5 / 10.4</b> | <b>0.913 / 0.897</b> |

**FIGURE 9.** Full PGA-UNet on the same BTXRD case (IMG000054) with a Gaussian-plateau prompt vs a binary box prompt, both at  $512 \times 512$ . The binary variant loses localization even though PSG and CAD are unchanged.

small head to directly regress the Dice; it collapsed to a near-constant output, with Spearman under 0.2 on both datasets, so it is not used.

Turned the other way around,  $Q$  can also rank 50 randomly sampled candidate boxes and keep only the best few (Table 16, Figure 11). For at least one of the top five boxes, 67% of BTXRD images and 82% of FracAtlas images have that box covering more than half the lesion, and the best of the five boxes covers 0.84 and 0.90 of the lesion area on average, respectively. To be clear, this is only a review-support tool:  $Q$  is not a true probability, and this shortlist does not replace the clinician's own review.

## VI. DISCUSSION

The question this article set out to answer was never simply whether a prompt is useful in general. The real question is: once a coarse box has already narrowed the search, does the way the model handles that prompt still matter? Attention U-Net without a box answers a different question altogether, and its low Dice, near 0.34 on FracAtlas, measures only how much of the task is finding the lesion. The comparisons that actually test the design are those where the same box is given to a plain Attention U-Net and to fine-tuned SAM-Med2D. In both cases, PGA-UNet keeps a margin: modest on BTXRD, clearer on FracAtlas, and largest on the Attention U-Net bottom-Dice group and on the smallest lesions. This is exactly the pattern one would expect if PSG and CAD are doing their job well: keeping the prompt active throughout the encoder and decoder, so a faint signal is not lost through downsampling. The single-split per-component ablation (Table 11) is consistent with this, with the Gaussian prompt consistently ahead under

**TABLE 12.** Monte Carlo cross-validation of PGA-UNet at  $512 \times 512$ , four repeated random image-level splits, mean  $\pm$  standard deviation.

| Dataset   | Prompt     | Dice              | IoU               | HD95           | CBL               |
|-----------|------------|-------------------|-------------------|----------------|-------------------|
| BTXRD     | Covering   | 0.769 $\pm$ 0.005 | 0.647 $\pm$ 0.006 | 25.0 $\pm$ 2.5 | 0.905 $\pm$ 0.005 |
|           | Off-center | 0.762 $\pm$ 0.004 | 0.639 $\pm$ 0.007 | 25.7 $\pm$ 2.0 | 0.901 $\pm$ 0.007 |
| FracAtlas | Covering   | 0.733 $\pm$ 0.007 | 0.595 $\pm$ 0.007 | 14.1 $\pm$ 2.7 | 0.909 $\pm$ 0.006 |
|           | Off-center | 0.728 $\pm$ 0.011 | 0.590 $\pm$ 0.010 | 15.6 $\pm$ 3.9 | 0.904 $\pm$ 0.013 |

**TABLE 13.** Efficiency in the tested environment (NVIDIA Tesla T4, batch size 1).

| Model         | Params  | GFLOPs | Size    | GPU ms | CPU ms |
|---------------|---------|--------|---------|--------|--------|
| PGA-UNet 256  | 2.96M   | 7.74   | 11.4 MB | 8.5    | 112    |
| PGA-UNet 512  | 2.96M   | 30.97  | 11.4 MB | 18.5   | 375    |
| SAM-Med2D 256 | 271.24M | 92.02  | 2443 MB | 50.1   | 1324   |


off-center prompts and on FracAtlas.

Beyond the main result, the study also yields two side findings. First, the loss comparison gives a fairly clear negative result: neither the size-conditioned Tversky loss nor the Focal Dice loss helps, and Focal Dice even makes things distinctly worse on FracAtlas, so we keep the default Dice + BCE objective. Second, the self-assessment score gives a moderate positive result:  $Q$  correlates with the true Dice at a rank level of about 0.5 to 0.7 without needing any label at all, and it succeeds where a small learned regression head does not, since that head ultimately just predicts a mean value. Even so,  $Q$  remains only a heuristic number, not yet calibrated, so the candidate list built from it should only be used to suggest where a clinician might look first, not to replace self-detection.

This study also has fairly concrete limitations. Inputs are resized to a fixed square, and the resolution results themselves show that this genuinely costs accuracy on the smallest lesions. The prompts in this article are all drawn from existing annotations, not directly drawn by a radiologist, and the protocol assumes the user has already found the correct region before drawing a box over it. Cases where the box only partially covers the lesion, is negative, or targets the wrong region are not addressed by this article. BTXRD and FracAtlas are trained completely separately, with no joint model and no cross-dataset test step, and patient-level separation cannot be confirmed from the public metadata. The Monte Carlo runs only show stability across splits, not superiority. The ablation experiment is currently still transitioning from a single split to multiple splits. In this round, SAM-Med2D remains the sole promptable foundation-model reference; comparison with newer promptable medical models is left for later. Finally, although the observed gains under prompt perturbation are fairly consistent with the original design intent, they have not yet been verified through direct feature-level analysis; so, once complete, the ablation experiment will only speak to the usefulness of the design as a whole, not yet confirm precisely which internal mechanism is playing the decisive role.

**TABLE 14.** Exploratory segmentation-loss comparison for PGA-UNet at  $512 \times 512$ , everything else fixed. Dice, IoU, CBL as covering / off-center. The default is not changed.

| Dataset   | Loss                     | Dice                 | IoU                  | HD95               | CBL                  |
|-----------|--------------------------|----------------------|----------------------|--------------------|----------------------|
| BTXRD     | Dice + BCE (default)     | <b>0.782 / 0.774</b> | <b>0.661 / 0.652</b> | 22.8 / 24.5        | <b>0.918 / 0.912</b> |
|           | Size-conditioned Tversky | 0.778 / 0.767        | 0.660 / 0.649        | 25.2 / 30.3        | 0.916 / 0.903        |
|           | Focal Dice               | 0.773 / 0.765        | 0.649 / 0.642        | 25.6 / 25.6        | 0.913 / 0.907        |
| FracAtlas | Dice + BCE (default)     | <b>0.727 / 0.713</b> | <b>0.585 / 0.570</b> | <b>10.5 / 10.4</b> | <b>0.913 / 0.897</b> |
|           | Size-conditioned Tversky | 0.709 / 0.680        | 0.561 / 0.530        | 11.0 / 12.2        | <b>0.918 / 0.888</b> |
|           | Focal Dice               | 0.642 / 0.578        | 0.493 / 0.425        | 105.5 / 110.1      | 0.802 / 0.732        |



Dice: 0.120 IoU: 0.064 Precision: 0.289 Recall: 0.076 HD95: 53.12 CBL: 0.621

**FIGURE 10.** Failure case on BTXRD (IMG000013, Dice 0.120): the prompt covers the lesion at the chest/rib margin, but PGA-UNet's prediction lands away from it. A covering box does not fix a case where the local image evidence is itself ambiguous.

## VII. CONCLUSION

This article presented PGA-UNet, a lightweight prompt-guided CNN for interactive bone X-ray lesion segmentation. The model's design turns a user-provided lesion-covering box into a dense spatial prior, then reuses this prior through PSG on the encoder side and CAD on the decoder side. On the BTXRD and FracAtlas datasets, evaluated under both controlled covering and off-center prompts, PGA-UNet consistently surpasses the conventional prompt-matched baselines as well as fine-tuned SAM-Med2D given the same box; this gap is clearest exactly where automatic localization fails and where lesions are smallest, while the model's parameter count is only about 1/92 that of SAM-Med2D. Repeated-split experiments show this trend is fairly stable, and a training-free self-assessment score also yields a ranking of prediction reliability that, while not yet calibrated, remains usable even without ground truth.

These conclusions are only valid within the protocol used: simulated boxes drawn around annotated lesions. They say nothing about unconstrained detection, off-target or partial-coverage boxes, calibrated confidence, patient-level generalization, or transfer to another dataset. In the near term, we will complete the multi-seed ablation experiment, extend the evaluation to more datasets and move toward a single shared model rather than training separately per dataset as at present, and add comparisons with more recent promptable medical models. In the longer term, we want to bring the method closer to real-world conditions: replacing the simulated boxes with looser, more realistic prompt protocols, removing the fixed-resolution resize step, and incorporating additional clinical information to improve both segmentation results and the quality of suggested regions. Finally, to reduce annotation cost and move closer to a tool usable in the clinic, the next direction is semi-supervised learning, exploiting the fact that drawing a coarse bounding box is much faster than drawing a

**TABLE 15.** Spearman correlation between each training-free cue and the true per-prompt Dice.  $Q = \text{mean}(S_{\text{sharp}}, S_{\text{stab}})$ .

| Cue                                    | BTXRD       |             | FracAtlas   |             |
|----------------------------------------|-------------|-------------|-------------|-------------|
|                                        | cover       | off-center  | cover       | off-center  |
| $S_{\text{sharp}}$ (mask decisiveness) | 0.42        | 0.39        | 0.47        | 0.43        |
| $S_{\text{stab}}$ (jitter stability)   | 0.64        | 0.56        | 0.63        | 0.45        |
| $Q$                                    | <b>0.70</b> | <b>0.60</b> | <b>0.65</b> | <b>0.48</b> |

**TABLE 16.** Candidate-region shortlist ranked by  $Q$ , keeping the 5 highest-scoring boxes out of 50 prompt-free boxes sampled per image. Coverage is the fraction of lesion pixels inside a box; full-coverage counts images whose best shortlisted box covers more than half the lesion.

| Dataset   | Mean cov. | Best-of-list cov. | Full-cov. rate |
|-----------|-----------|-------------------|----------------|
| BTXRD     | 0.52      | 0.84              | 66.8%          |
| FracAtlas | 0.62      | 0.90              | 81.9%          |

pixel-accurate mask. We want to make use of a large pool of data that only has cheap box labels, combined with a small amount of accurate masks, instead of requiring a detailed mask for every sample as at present, together with turning the self-assessment score into a calibrated signal usable for suggesting suspicious regions in an actual clinical review workflow.

## DATA AVAILABILITY

BTXRD and FracAtlas are public datasets cited in this article. The manuscript reports experiments on image-level train/validation/test splits derived from those datasets. Code and split details should be released with the final reproducibility package.

## AUTHOR CONTRIBUTIONS

Thong Luc: conceptualization, methodology, software, investigation, formal analysis, visualization, writing of the original draft. Nguyen Huu Binh: software support, data preparation, experiment verification, review and editing. Ly Quoc Ngoc: supervision, validation, review and editing.

## ETHICS STATEMENT


This study used public de-identified radiograph datasets and did not involve prospective data collection, patient contact, or intervention.

## ACKNOWLEDGMENT

This research is funded by Vietnam National University Ho Chi Minh City (VNU-HCM) under Grant number A2025-18-03.

OpenAI ChatGPT was used only for vocabulary and grammar refinement in selected non-technical manuscript text. All technical content, experimental design, mathematical formulation, figures, tables, and conclusions were prepared and verified by the authors, who take full responsibility for the final manuscript.





**FIGURE 11.** Candidate-region shortlist on a BTXRD case. Fifty boxes are sampled with no prompt, scored by the self-assessment score  $Q$ , and the five highest-scoring boxes are shown for the clinician to accept, adjust, or reject. A FracAtlas case is shown in the supplementary material.

## REFERENCES

- [1] A. Kirillov *et al.*, "Segment Anything," in *Proc. IEEE/CVF Int. Conf. Comput. Vis.*, 2023, pp. 4015–4026.
- [2] J. Cheng *et al.*, "SAM-Med2D," *arXiv preprint arXiv:2308.16184*, 2023.
- [3] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation," in *MICCAI*, 2015, pp. 234–241.
- [4] O. Oktay *et al.*, "Attention U-Net: Learning Where to Look for the Pancreas," in *Medical Imaging with Deep Learning*, 2018.
- [5] N. Xu, B. Price, S. Cohen, J. Yang, and T. S. Huang, "Deep Interactive Object Selection," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit.*, 2016, pp. 373–381.
- [6] G. Wang *et al.*, "DeepIGeoS: A Deep Interactive Geodesic Framework for Medical Image Segmentation," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 41, no. 7, pp. 1559–1572, 2019.
- [7] S. Yao *et al.*, "A Radiograph Dataset for the Classification, Localization, and Segmentation of Primary Bone Tumors," *Scientific Data*, vol. 12, no. 1, p. 88, 2025.
- [8] I. Abdeen *et al.*, "FracAtlas: A Dataset for Fracture Classification, Localization and Segmentation of Musculoskeletal Radiographs," *Scientific Data*, vol. 10, no. 1, p. 521, 2023.
- [9] J. Ma *et al.*, "Segment Anything in Medical Images," *Nat. Commun.*, vol. 15, no. 1, p. 654, 2024.
- [10] F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, "nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation," *Nat. Methods*, vol. 18, pp. 203–211, 2021.
- [11] H. E. Wong, M. Rakic, J. Guttig, and A. V. Dalca, "ScribblePrompt: Fast and Flexible Interactive Segmentation for Any Biomedical Image," in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2024, pp. 207–229.
- [12] G. Dong *et al.*, "An efficient segment anything model for the segmentation of medical images," *Sci. Rep.*, vol. 14, art. 19425, 2024.
- [13] J. Wu *et al.*, "Medical SAM adapter: Adapting segment anything model for medical image segmentation," *Med. Image Anal.*, vol. 102, art. 103547, 2025.

**NGUYEN HUU BINH** Nguyen Huu Binh is currently an undergraduate student in Computer Science with the Faculty of Information Technology, University of Science, Vietnam National University Ho Chi Minh City, Vietnam. His research interests include computer vision, deep learning, and image segmentation. His current research focuses on deep learning methods for medical image segmentation using visual prompts.



**LY QUOC NGOC** Ly Quoc Ngoc (Associate Professor, Ph.D.) is Head of the Department of Computer Vision and Cognitive Cybernetics, Faculty of Information Technology, University of Science, Vietnam National University Ho Chi Minh City, Vietnam. He received the B.Sc. degree in Mechanics and Mathematics, the M.Sc. degree in Computer Science, and the D.Sc. degree in Computer Science, all from the University of Science, Ho Chi Minh City, in 1987, 1996, and 2009, respectively.

His research interests include artificial intelligence, machine learning, computer vision, digital image and video processing, computer graphics, multivariate statistical analysis, and medical image analysis. He has authored or co-authored nearly 70 papers in conference proceedings and academic journals, and has served as principal investigator of several research projects in intelligent vision systems.



**THONG LUC** Thong Luc is a final-year undergraduate student in the B.Sc. program in Computer Science with the Faculty of Information Technology, University of Science, Vietnam National University Ho Chi Minh City, Vietnam. His research interests include computer vision, deep learning, medical image analysis, image segmentation, and interactive prompt-guided segmentation. His current work focuses on prompt-guided segmentation for bone radiographic images using deep learning.